Slick Trace investigation report — case_0007

Generated 2026-09-14T07:36:12.346264+00:00 · reproducible investigation report

What this document is not. This is not a court-ready dossier, a legal finding, or an accusation. It records what the system computed and under which assumptions.

Finding

Leading candidate: 419240128. This is a ranking by physical consistency under tested assumptions — an investigative lead for a human analyst, not proof of responsibility.

Scene

Field	Value
Incident	case_0007
Acquired (UTC)	2026-09-10T04:30:00+00:00
Sensor	synthetic SAR-like (VV proxy, single channel)
Data status	synthetic
Pixel spacing	433.59 m
Label source	generator ground truth (synthetic) — not a human or independently confirmed label

Evidence layers

Layer	Provenance
AIS	synthetic (fictional MMSIs; no real broadcast data)
Forcing	slicktrace_core.synth_generator constant+spatial field v0.2
Licence	Synthetic data generated by this repository; released under the repository licence. No third-party data terms apply.

Traffic filtering

State	Vessels
candidate	4
excluded	2

Unassessed is reported separately from excluded: missing evidence is not exoneration, and collapsing the two would silently convert one into the other.

Rank fragility

Leader retained in 1.0 of 80 tested perturbations. No tested perturbation changed the leader.

Absence of a flip is evidence of stability across the tested set, not a proof of stability: single-axis and two-axis perturbations only.

File integrity

File	SHA-256 (first 16)
attribution_result.json	d4197c23e8301954
coverage_result.json	dc3f4b9c047730e4
forcing.npz	f8958ffe32637fe5
forecast_result.json	a6eb9bdff5f46e88
fragility_result.json	76f3fafe3bc39f2a
historical_mapping.json	41bb0c366861beee
manifest.json	4d697f3633004fce
mask_groundtruth.png	a028649edc062bf9
naive_baseline_result.json	ac1e534b2baf6b20
next_observation_result.json	7e46bac46b21f159
scene.png	b8d9917418fd516a

Hashes verify file integrity only. They do not establish that an observation is true, that a label is correct, or that this analysis is legally sufficient.

Uncertainty

• Rankings express physical consistency under tested assumptions. They are not proof of responsibility.

• Ensemble and perturbation-grid win frequencies are evidence scores, not calibrated probabilities. No calibration against held-out real outcomes exists.

• Where no candidate clears the absolute fit gate the system reports unknown source rather than naming the least-bad option.

• Synthetic ranking accuracy is not a measurement of real-world attribution accuracy.

Restated. This is not a court-ready dossier, a legal finding, or an accusation. It records what the system computed and under which assumptions.